

```

1
2 # %%%%%%%%% A. NOTEBOOK SETUP %%%%%%%%%
3 # This cell sets up the environment by importing the necessary libraries.
4
5 import torch
6 import numpy as np
7 import pandas as pd
8 import networkx as nx
9 import matplotlib.pyplot as plt
10 import plotly.graph_objects as go
11 from plotly.subplots import make_subplots
12 from IPython.display import display, Markdown
13
14 # Set a consistent style for plots
15 plt.style.use('seaborn-v0_8-whitegrid')
16
17 print("Libraries imported successfully.")
18
19
20 # %%%%%%%%% B. SIMULATION CORE: GRAPH GENERATION %%%%%%%%%
21 # This cell contains functions to generate the different task dependency graphs
22 # that will be used in our assessment scenarios.
23
24 def generate_chain_graph(num_nodes):
25     """Generates a simple chain graph (0 -> 1 -> 2 -> ...)."""
26     G = nx.DiGraph()
27     G.add_nodes_from(range(num_nodes))
28     for i in range(num_nodes - 1):
29         G.add_edge(i, i + 1)
30     return G
31
32 def generate_tree_graph(num_nodes, fan_out=3):
33     """Generates a tree graph with one root branching out."""
34     if num_nodes == 0:
35         return nx.DiGraph()
36     G = nx.DiGraph()
37     G.add_node(0) # Root node
38     nodes_to_process = [0]
39     next_node_id = 1
40     while nodes_to_process and next_node_id < num_nodes:
41         current_node = nodes_to_process.pop(0)
42         for i in range(fan_out):
43             if next_node_id < num_nodes:
44                 G.add_node(next_node_id)
45                 G.add_edge(current_node, next_node_id)
46                 nodes_to_process.append(next_node_id)
47                 next_node_id += 1
48     return G
49
50 def generate_inverted_tree_graph(num_nodes, fan_in=3):
51     """Generates an inverted tree with many roots converging to one final task."""
52     if num_nodes == 0:
53         return nx.DiGraph()
54     G = nx.DiGraph()
55     final_task = num_nodes - 1
56     G.add_node(final_task)
57     # The last node is the sink. All other nodes are potential sources.
58     # We will form layers that converge on the final task.
59
60     # This is a simplified generation logic
61     num_leaves = num_nodes - 1
62     for i in range(num_leaves):
63         G.add_node(i)
64         G.add_edge(i, final_task)
65
66     return G
67
68 def generate_dag(num_nodes, sparsity=0.2):
69     """Generates a random Directed Acyclic Graph (DAG)."""
70     G = nx.DiGraph()
71     G.add_nodes_from(range(num_nodes))
72     for i in range(num_nodes):
73         for j in range(i + 1, num_nodes):
74             if np.random.rand() < sparsity:

```

```

75         G.add_edge(i, j)
76     return G
77
78     print("Graph generation functions are defined.")
79
80
81     # %%%%%%%%% C. SIMULATION CORE: THE SIMULATOR CLASS %%%%%%%%%
82     # This is the heart of the benchmark. The `CurriculumSimulator` class runs the
83     # simulation for a given curriculum, task graph, and set of dynamic parameters.
84     # It now includes logic for all the new curricula we designed.
85
86     class CurriculumSimulator:
87         """Runs a curriculum learning simulation for a given configuration."""
88
89         def __init__(self, config, graph, curriculum_type):
90             self.config = config
91             self.G = graph
92             self.curriculum = curriculum_type
93
94             # Pre-compute graph properties
95             self.adj = [list(self.G.successors(i)) for i in range(self.config.NUM_TASKS)]
96             self.parents = [list(self.G.predecessors(i)) for i in range(self.config.NUM_TASKS)]
97             try:
98                 self.topological_order = list(nx.topological_sort(self.G))
99                 self.reverse_topological_order = self.topological_order[::-1]
100             except nx.NetworkXUnfeasible: # For graphs with no clear topological sort (e.g., disconnected)
101                 self.topological_order = list(range(self.config.NUM_TASKS))
102                 self.reverse_topological_order = list(range(self.config.NUM_TASKS))[::-1]
103
104
105             # Initialize performance and history trackers
106             self.P = torch.full((self.config.NUM_TASKS,), 0.01)
107             self.P_history = [self.P.clone()]
108             self.sampling_probs_history = [torch.full((self.config.NUM_TASKS,), 1.0 / self.config.NUM_TASKS)]
109
110             # State for specific curricula
111             self.fast_ema = torch.zeros(self.config.NUM_TASKS)
112             self.slow_ema = torch.zeros(self.config.NUM_TASKS)
113             self.perf_window = torch.zeros((self.config.VARIANCE_WINDOW, self.config.NUM_TASKS))
114             self.topological_idx = 0
115
116
117         def _get_sampling_probs(self, epoch):
118             """Calculates sampling probabilities based on the chosen curriculum."""
119             # --- BASELINES AND HEURISTICS ---
120             if self.curriculum == 'random' or epoch == 0:
121                 return torch.full((self.config.NUM_TASKS,), 1.0 / self.config.NUM_TASKS)
122
123             if self.curriculum == 'topological':
124                 probs = torch.zeros(self.config.NUM_TASKS)
125                 current_task = self.topological_order[self.topological_idx]
126                 probs[current_task] = 1.0
127                 # Move to the next task if current one is mastered
128                 if self.P[current_task] >= self.config.PERFORMANCE_THRESHOLD:
129                     self.topological_idx = min(self.topological_idx + 1, self.config.NUM_TASKS - 1)
130                 return probs
131
132             if self.curriculum == 'reverse_topological':
133                 probs = torch.zeros(self.config.NUM_TASKS)
134                 current_task = self.reverse_topological_order[self.topological_idx]
135                 probs[current_task] = 1.0
136                 if self.P[current_task] >= self.config.PERFORMANCE_THRESHOLD:
137                     self.topological_idx = min(self.topological_idx + 1, self.config.NUM_TASKS - 1)
138                 return probs
139
140             if self.curriculum == 'oracle':
141                 # This heuristic Oracle identifies tasks whose parents are all mastered
142                 # and samples uniformly from that set of "unlocked" tasks.
143                 unlocked_tasks = []
144                 for i in range(self.config.NUM_TASKS):
145                     if self.P[i] < self.config.PERFORMANCE_THRESHOLD: # Task is not yet mastered
146                         parents_mastered = all(self.P[p] >= self.config.PERFORMANCE_THRESHOLD for p in self.parents[i])
147                         if parents_mastered:
148                             unlocked_tasks.append(i)
149

```

```

150         if not unlocked_tasks: # If nothing is unlocked, fall back to random
151             return torch.full((self.config.NUM_TASKS,), 1.0 / self.config.NUM_TASKS)
152
153         probs = torch.zeros(self.config.NUM_TASKS)
154         for task_idx in unlocked_tasks:
155             probs[task_idx] = 1.0
156         return probs / probs.sum()
157
158     # --- ADAPTIVE CURRICULA ---
159     last_P = self.P_history[-1]
160
161     if self.curriculum == 'lp_finite_diff':
162         second_last_P = self.P_history[-2] if len(self.P_history) > 1 else last_P
163         progress = torch.abs(last_P - second_last_P)
164
165     elif self.curriculum == 'lp_ema_diff':
166         self.fast_ema = self.config.FAST_EMA_ALPHA * last_P + (1 - self.config.FAST_EMA_ALPHA) * self.fast_ema
167         self.slow_ema = self.config.SLOW_EMA_ALPHA * last_P + (1 - self.config.SLOW_EMA_ALPHA) * self.slow_ema
168         progress = torch.abs(self.fast_ema - self.slow_ema)
169
170     elif self.curriculum == 'variance':
171         # Update performance window
172         self.perf_window = torch.roll(self.perf_window, shifts=-1, dims=0)
173         self.perf_window[-1, :] = last_P
174         # Progress is the variance over the last K epochs
175         progress = torch.var(self.perf_window, dim=0)
176
177     else:
178         raise ValueError(f"Unknown curriculum: {self.curriculum}")
179
180     # Convert progress scores to probabilities using softmax
181     if torch.all(progress == 0): # Avoid NaN if progress is zero everywhere
182         return torch.full((self.config.NUM_TASKS,), 1.0 / self.config.NUM_TASKS)
183
184     return torch.nn.functional.softmax(progress / self.config.SOFTMAX_TEMP, dim=0)
185
186
187 def run(self):
188     """Executes the full simulation loop."""
189     for epoch in range(self.config.NUM_EPOCHS):
190         sampling_probs = self._get_sampling_probs(epoch)
191         self.sampling_probs_history.append(sampling_probs)
192
193         # Get sample counts for each task
194         S = torch.multinomial(sampling_probs, self.config.TOTAL_SAMPLES_PER_EPOCH, replacement=True)
195         S_counts = torch.bincount(S, minlength=self.config.NUM_TASKS).float()
196
197         # Calculate performance change (P_dot)
198         current_P = self.P
199         P_dot = torch.zeros(self.config.NUM_TASKS)
200
201         for i in range(self.config.NUM_TASKS):
202             parent_contribution = sum(S_counts[p_idx] for p_idx in self.parents[i])
203             total_stimulus = S_counts[i] + self.config.GAMMA * parent_contribution
204
205             children_gate = 1.0
206             if self.adj[i]: # If task i has children
207                 children_gate = torch.prod(torch.tensor([current_P[c_idx] for c_idx in self.adj[i]]))
208
209             growth = total_stimulus * children_gate * (1 - current_P[i])
210             forgetting = self.config.LAMBDA * current_P[i]
211
212             # Normalize by total samples to make update step size independent of sample count
213             P_dot[i] = (growth - forgetting) / self.config.TOTAL_SAMPLES_PER_EPOCH
214
215         # Update performance (Euler integration with dt=1)
216         new_P = current_P + P_dot
217         self.P = torch.clamp(new_P, 0, 1)
218         self.P_history.append(self.P.clone())
219
220         # Early stopping if threshold is met
221         if torch.all(self.P >= self.config.PERFORMANCE_THRESHOLD):
222             # Pad history to full length for consistent analysis
223             pad_len = self.config.NUM_EPOCHS - epoch - 1
224             if pad_len > 0:

```

```

225         self.P_history.extend([self.P.clone()] * pad_len)
226     break
227
228     def calculate_metrics(self):
229         """Calculates summary metrics after the simulation."""
230         P_history_tensor = torch.stack(self.P_history)
231         mean_perf_per_epoch = [p.mean().item() for p in self.P_history]
232
233         # Learning Efficiency
234         efficiency = sum(mean_perf_per_epoch)
235
236         # Time to Thresholds
237         time_to_threshold = -1
238         time_to_first_mastery = -1
239
240         for i, p_epoch in enumerate(self.P_history):
241             if time_to_first_mastery == -1 and torch.any(p_epoch >= self.config.PERFORMANCE_THRESHOLD):
242                 time_to_first_mastery = i
243             if time_to_threshold == -1 and torch.all(p_epoch >= self.config.PERFORMANCE_THRESHOLD):
244                 time_to_threshold = i
245
246         # Final Performance Variance
247         final_perf_variance = torch.var(self.P_history[-1]).item()
248
249         return {
250             'efficiency': efficiency,
251             'time_to_threshold': time_to_threshold,
252             'time_to_first_mastery': time_to_first_mastery,
253             'final_perf_variance': final_perf_variance,
254         }
255
256     print("CurriculumSimulator class defined.")
257
258
259     # %%%%%%%%% D. BENCHMARKING FRAMEWORK %%%%%%%%%
260     # This cell defines the full benchmark. It sets up the scenarios and curricula
261     # to test, then runs the sweep, storing results in a pandas DataFrame.
262
263     class Config:
264         # Graph params
265         NUM_TASKS = 12
266         SPARSITY = 0.3
267         # Dynamics params
268         GAMMA = 0.5
269         LAMBDA = 0.01
270         # Simulation params
271         NUM_EPOCHS = 200
272         TOTAL_SAMPLES_PER_EPOCH = 100
273         PERFORMANCE_THRESHOLD = 0.9
274         # LP/Variance Curriculum params
275         SOFTMAX_TEMP = 0.1
276         FAST_EMA_ALPHA = 0.3
277         SLOW_EMA_ALPHA = 0.05
278         VARIANCE_WINDOW = 5 # For variance curriculum
279
280     # --- Define Scenarios ---
281     scenarios = {
282         "1_Chain_Simple": {
283             "graph_type": "chain",
284             "LAMBDA": 0.0,
285             "GAMMA": 0.5
286         },
287         "2_Chain_HighForget": {
288             "graph_type": "chain",
289             "LAMBDA": 0.05, # High forgetting
290             "GAMMA": 0.5
291         },
292         "3_Tree_Divergent": {
293             "graph_type": "tree",
294             "LAMBDA": 0.01,
295             "GAMMA": 0.5
296         },
297         "4_InvertedTree_Convergent": {
298             "graph_type": "inverted_tree",
299             "LAMBDA": 0.01,

```

```

300     "GAMMA": 0.5
301 },
302 "5_ComplexDAG_LowTransfer": {
303     "graph_type": "dag",
304     "LAMBDA": 0.01,
305     "GAMMA": 0.1 # Low knowledge transfer
306 },
307 }
308
309 # --- Define Curricula ---
310 curricula_to_test = [
311     "random",
312     "topological",
313     "reverse_topological",
314     "lp_finite_diff",
315     "lp_ema_diff",
316     "variance",
317     "oracle"
318 ]
319
320 def run_benchmark_sweep():
321     """Runs the full experimental sweep over all scenarios and curricula."""
322     results = []
323
324     # Generate a fixed set of graphs for consistency across scenarios of the same type
325     base_graphs = {
326         "chain": generate_chain_graph(Config.NUM_TASKS),
327         "tree": generate_tree_graph(Config.NUM_TASKS),
328         "inverted_tree": generate_inverted_tree_graph(Config.NUM_TASKS),
329         "dag": generate_dag(Config.NUM_TASKS, Config.SPARSITY)
330     }
331
332     total_runs = len(scenarios) * len(curricula_to_test)
333     run_count = 0
334
335     for s_name, s_params in scenarios.items():
336         for curriculum in curricula_to_test:
337             run_count += 1
338             print(f"Running ({run_count}/{total_runs}): Scenario='{s_name}', Curriculum='{curriculum}'...")
339
340             # Setup config for this run
341             config = Config()
342             config.LAMBDA = s_params["LAMBDA"]
343             config.GAMMA = s_params["GAMMA"]
344             graph = base_graphs[s_params["graph_type"]]
345
346             # Run simulation
347             simulator = CurriculumSimulator(config, graph, curriculum)
348             simulator.run()
349             metrics = simulator.calculate_metrics()
350
351             # Store results
352             result_row = {
353                 "scenario": s_name,
354                 "curriculum": curriculum,
355                 "graph_type": s_params["graph_type"],
356                 **metrics
357             }
358             results.append(result_row)
359
360     print("\nBenchmark sweep complete.")
361     return pd.DataFrame(results)
362
363 # --- Execute the sweep ---
364 results_df = run_benchmark_sweep()
365 display(results_df)
366
367
368 # %%%%%%%%% E. REGRET CALCULATION %%%%%%%%%
369 # This cell processes the raw results DataFrame to calculate regret metrics.
370 # Regret is calculated by comparing each curriculum's performance to the
371 # 'oracle' in the same scenario.
372
373 def calculate_regret(df):
374     """Calculates regret metrics based on the oracle's performance."""

```

```

375 oracle_metrics = df[df['curriculum'] == 'oracle'].set_index('scenario')
376
377 regret_data = []
378
379 for index, row in df.iterrows():
380     if row['curriculum'] == 'oracle':
381         # Oracle has zero regret by definition
382         efficiency_regret = 0
383         time_regret = 0
384     else:
385         scenario = row['scenario']
386         oracle_row = oracle_metrics.loc[scenario]
387
388         # Efficiency Regret
389         efficiency_regret = oracle_row['efficiency'] - row['efficiency']
390
391         # Time Regret
392         # If oracle or curriculum failed (-1), regret is complex.
393         # We'll define it as a large number if the curriculum fails but oracle succeeds.
394         if row['time_to_threshold'] == -1 and oracle_row['time_to_threshold'] != -1:
395             time_regret = Config.NUM_EPOCHS # Max penalty
396         elif row['time_to_threshold'] == -1 and oracle_row['time_to_threshold'] == -1:
397             time_regret = 0 # Both failed
398         elif row['time_to_threshold'] != -1 and oracle_row['time_to_threshold'] == -1:
399             time_regret = -Config.NUM_EPOCHS # Actually did better than failing oracle!
400         else:
401             time_regret = row['time_to_threshold'] - oracle_row['time_to_threshold']
402
403     regret_data.append({
404         'efficiency_regret': efficiency_regret,
405         'time_regret': time_regret
406     })
407
408 regret_df = pd.DataFrame(regret_data, index=df.index)
409 return df.join(regret_df)
410
411 # --- Execute regret calculation ---
412 results_with_regret_df = calculate_regret(results_df)
413 display(Markdown("### Results DataFrame with Regret Metrics"))
414 display(results_with_regret_df)
415
416
417 # %%%%%%%%% F. VISUALIZATION SUITE %%%%%%%%%
418 # This final part generates the visualizations. It includes functions for both
419 # the detailed bar charts and the high-level summary star plots, which provide
420 # a compelling final comparison.
421
422 def plot_bar_charts(df, scenario_name):
423     """Generates a bar chart comparing all curricula for a given scenario."""
424
425     scenario_df = df[df['scenario'] == scenario_name].set_index('curriculum')
426
427     metrics_to_plot = [
428         'efficiency', 'time_to_threshold',
429         'efficiency_regret', 'time_regret',
430         'time_to_first_mastery', 'final_perf_variance'
431     ]
432
433     fig, axes = plt.subplots(3, 2, figsize=(18, 15))
434     axes = axes.ravel()
435     fig.suptitle(f"Metric Comparison for Scenario: {scenario_name}", fontsize=20)
436
437     for i, metric in enumerate(metrics_to_plot):
438         data = scenario_df[metric].sort_values(ascending=False)
439         colors = plt.cm.viridis(np.linspace(0, 1, len(data)))
440         bars = axes[i].bar(data.index, data.values, color=colors)
441         axes[i].set_title(metric.replace('_', ' ').title(), fontsize=14)
442         axes[i].tick_params(axis='x', rotation=45)
443         axes[i].bar_label(bars, fmt='%.1f')
444
445     plt.tight_layout(rect=[0, 0, 1, 0.96])
446     plt.show()
447
448
449 def plot_star_plots(df):
450     """

```

```

451 Generates summary star plots (radar charts) comparing curricula across all scenarios.
452 This is the final, compelling visualization of strengths and weaknesses.
453 """
454
455 # --- 1. Normalize Metrics ---
456 # Star plots require metrics to be on a similar scale, where 'bigger is better'.
457 df_norm = df.copy()
458
459 # For 'lower is better' metrics, we invert them.
460 for col in ['time_to_threshold', 'time_to_first_mastery', 'final_perf_variance', 'efficiency_regret', 'time_regret']:
461     # Handle -1 (failure) case for time metrics
462     if 'time' in col:
463         max_val = Config.NUM_EPOCHS
464         # Replace -1 with max penalty
465         df_norm[col] = df_norm[col].replace(-1, max_val * 1.1)
466
467     # Invert the metric so higher is better
468     # Add small epsilon to avoid division by zero
469     min_val = df_norm[col].min()
470     max_val = df_norm[col].max()
471     if max_val - min_val == 0:
472         df_norm[f'norm_{col}'] = 0.5
473     else:
474         df_norm[f'norm_{col}'] = (max_val - df_norm[col]) / (max_val - min_val)
475
476 # For 'higher is better' metrics, we just scale them from 0 to 1.
477 for col in ['efficiency']:
478     min_val = df_norm[col].min()
479     max_val = df_norm[col].max()
480     if max_val - min_val == 0:
481         df_norm[f'norm_{col}'] = 0.5
482     else:
483         df_norm[f'norm_{col}'] = (df_norm[col] - min_val) / (max_val - min_val)
484
485 # --- 2. Create Plots ---
486 metrics_for_star = [
487     'norm_efficiency', 'norm_time_to_threshold',
488     'norm_time_to_first_mastery', 'norm_final_perf_variance',
489     'norm_efficiency_regret', 'norm_time_regret'
490 ]
491
492 # Clean labels for the plot
493 theta_labels = [m.replace('norm_', '').replace('_', ' ').replace('Regret', ' (Low Regret)').title() for m in metrics_for_star]
494
495 # Get unique scenarios and curricula
496 scenarios = df_norm['scenario'].unique()
497 curricula = df_norm[df_norm['curriculum'] != 'oracle']['curriculum'].unique() # Exclude oracle from comparison
498
499 fig = make_subplots(
500     rows=1, cols=len(scenarios),
501     specs=[['type': 'polar']] * len(scenarios),
502     subplot_titles=scenarios
503 )
504
505 for i, scenario in enumerate(scenarios):
506     scenario_df = df_norm[df_norm['scenario'] == scenario]
507
508     for curriculum in curricula:
509         row = scenario_df[scenario_df['curriculum'] == curriculum]
510         if not row.empty:
511             r_values = row[metrics_for_star].values.flatten().tolist()
512             fig.add_trace(
513                 go.Scatterpolar(
514                     r=r_values,
515                     theta=theta_labels,
516                     fill='toself',
517                     name=curriculum,
518                     legendgroup=curriculum,
519                     showlegend=(i==0) # Show legend only for the first subplot
520                 ),
521                 row=1, col=i+1
522             )
523
524 fig.update_layout(
525     height=600.

```

```

526         width=400*len(scenarios),
527         title_text="Curriculum Performance Profiles (Normalized)",
528         legend_title_text='Curricula',
529         polar=dict(radialaxis=dict(visible=True, range=[0, 1]))
530     )
531     fig.show()
532
533
534 # --- Execute Visualizations ---
535 display(Markdown("## Detailed Bar Chart Comparisons"))
536 display(Markdown("These charts provide a detailed look at curriculum performance in each specific scenario."))
537 for s_name in results_with_regret_df['scenario'].unique():
538     plot_bar_charts(results_with_regret_df, s_name)
539
540 display(Markdown("## Global Star Plot Comparison"))
541 display(Markdown("""
542 This star plot is the final, high-level summary. Each point on the star represents a normalized performance metric,
543 where **a larger area is universally better**. This visualization quickly reveals the unique strengths and weaknesses
544 of each curriculum across the different environmental challenges. For example, a curriculum that excels in the 'High F
545 scenario will have a large area in that respective plot.
546 """))
547 plot_star_plots(results_with_regret_df)

```



```

↳ Libraries imported successfully.
Graph generation functions are defined.
CurriculumSimulator class defined.
Running (1/35): Scenario='1_Chain_Simple', Curriculum='random'...
Running (2/35): Scenario='1_Chain_Simple', Curriculum='topological'...
Running (3/35): Scenario='1_Chain_Simple', Curriculum='reverse_topological'...
Running (4/35): Scenario='1_Chain_Simple', Curriculum='lp_finite_diff'...
Running (5/35): Scenario='1_Chain_Simple', Curriculum='lp_ema_diff'...
Running (6/35): Scenario='1_Chain_Simple', Curriculum='variance'...
Running (7/35): Scenario='1_Chain_Simple', Curriculum='oracle'...
Running (8/35): Scenario='2_Chain_HighForget', Curriculum='random'...
Running (9/35): Scenario='2_Chain_HighForget', Curriculum='topological'...
Running (10/35): Scenario='2_Chain_HighForget', Curriculum='reverse_topological'...
Running (11/35): Scenario='2_Chain_HighForget', Curriculum='lp_finite_diff'...
Running (12/35): Scenario='2_Chain_HighForget', Curriculum='lp_ema_diff'...
Running (13/35): Scenario='2_Chain_HighForget', Curriculum='variance'...
Running (14/35): Scenario='2_Chain_HighForget', Curriculum='oracle'...
Running (15/35): Scenario='3_Tree_Divergent', Curriculum='random'...
Running (16/35): Scenario='3_Tree_Divergent', Curriculum='topological'...
Running (17/35): Scenario='3_Tree_Divergent', Curriculum='reverse_topological'...
Running (18/35): Scenario='3_Tree_Divergent', Curriculum='lp_finite_diff'...
Running (19/35): Scenario='3_Tree_Divergent', Curriculum='lp_ema_diff'...
Running (20/35): Scenario='3_Tree_Divergent', Curriculum='variance'...
Running (21/35): Scenario='3_Tree_Divergent', Curriculum='oracle'...
Running (22/35): Scenario='4_InvertedTree_Convergent', Curriculum='random'...
Running (23/35): Scenario='4_InvertedTree_Convergent', Curriculum='topological'...
Running (24/35): Scenario='4_InvertedTree_Convergent', Curriculum='reverse_topological'...
Running (25/35): Scenario='4_InvertedTree_Convergent', Curriculum='lp_finite_diff'...
Running (26/35): Scenario='4_InvertedTree_Convergent', Curriculum='lp_ema_diff'...
Running (27/35): Scenario='4_InvertedTree_Convergent', Curriculum='variance'...
Running (28/35): Scenario='4_InvertedTree_Convergent', Curriculum='oracle'...
Running (29/35): Scenario='5_ComplexDAG_LowTransfer', Curriculum='random'...
Running (30/35): Scenario='5_ComplexDAG_LowTransfer', Curriculum='topological'...
Running (31/35): Scenario='5_ComplexDAG_LowTransfer', Curriculum='reverse_topological'...
Running (32/35): Scenario='5_ComplexDAG_LowTransfer', Curriculum='lp_finite_diff'...
Running (33/35): Scenario='5_ComplexDAG_LowTransfer', Curriculum='lp_ema_diff'...
Running (34/35): Scenario='5_ComplexDAG_LowTransfer', Curriculum='variance'...
Running (35/35): Scenario='5_ComplexDAG_LowTransfer', Curriculum='oracle'...

```

Benchmark sweep complete.

	scenario	curriculum	graph_type	efficiency	time_to_threshold	time_to_first_mastery	final_perf_
0	1_Chain_Simple	random	chain	170.333115	66	17	6.8i
1	1_Chain_Simple	topological	chain	59.916556	-1	29	1.9
2	1_Chain_Simple	reverse_topological	chain	188.156021	24	2	0.00
3	1_Chain_Simple	lp_finite_diff	chain	172.032845	64	5	6.9i
4	1_Chain_Simple	lp_ema_diff	chain	111.270636	184	6	1.3
5	1_Chain_Simple	variance	chain	169.361163	67	16	7.4
6	1_Chain_Simple	oracle	chain	61.887785	-1	29	1.6i
7	2_Chain_HighForget	random	chain	167.824185	63	20	6.6i
8	2_Chain_HighForget	topological	chain	58.821264	-1	28	1.7i
9	2_Chain_HighForget	reverse_topological	chain	187.235618	24	2	1.1i
10	2_Chain_HighForget	lp_finite_diff	chain	172.371547	65	6	7.1i
11	2_Chain_HighForget	lp_ema_diff	chain	109.621859	185	6	2.1i
12	2_Chain_HighForget	variance	chain	169.088831	67	16	6.7i
13	2_Chain_HighForget	oracle	chain	60.436184	-1	29	1.9i
14	3_Tree_Divergent	random	tree	186.776147	55	17	7.6i
15	3_Tree_Divergent	topological	tree	38.272161	-1	200	8.6i
16	3_Tree_Divergent	reverse_topological	tree	188.741723	24	2	8.5i
17	3_Tree_Divergent	lp_finite_diff	tree	186.622626	47	13	7.1i
18	3_Tree_Divergent	lp_ema_diff	tree	177.935767	81	10	7.8i
19	3_Tree_Divergent	variance	tree	186.499868	53	16	6.7i
20	3_Tree_Divergent	oracle	tree	31.211558	-1	-1	5.5i
21	4_InvertedTree_Convergent	random	inverted_tree	177.925998	32	3	5.4i

22	4_InvertedTree_Convergent	topological	inverted_tree	188.712836	23	3	1.1%
23	4_InvertedTree_Convergent	reverse_topological	inverted_tree	188.009549	24	2	4.2%
24	4_InvertedTree_Convergent	lp_finite_diff	inverted_tree	177.705068	33	2	6.4%
25	4_InvertedTree_Convergent	lp_ema_diff	inverted_tree	169.647550	62	3	8.9%
26	4_InvertedTree_Convergent	variance	inverted_tree	176.611744	32	3	6.6%
27	4_InvertedTree_Convergent	oracle	inverted_tree	175.801671	27	4	7.2%
28	5_ComplexDAG_LowTransfer	random	dag	171.140873	79	16	7.6%
29	5_ComplexDAG_LowTransfer	topological	dag	20.379538	-1	-1	6.4%
30	5_ComplexDAG_LowTransfer	reverse_topological	dag	187.979976	24	2	6.6%
31	5_ComplexDAG_LowTransfer	lp_finite_diff	dag	175.494895	65	6	9.3%
32	5_ComplexDAG_LowTransfer	lp_ema_diff	dag	144.764146	138	6	1.0%
33	5_ComplexDAG_LowTransfer	variance	dag	171.462537	78	17	7.5%
34	5_ComplexDAG_LowTransfer	oracle	dag	42.932153	-1	43	1.5%

Results DataFrame with Regret Metrics

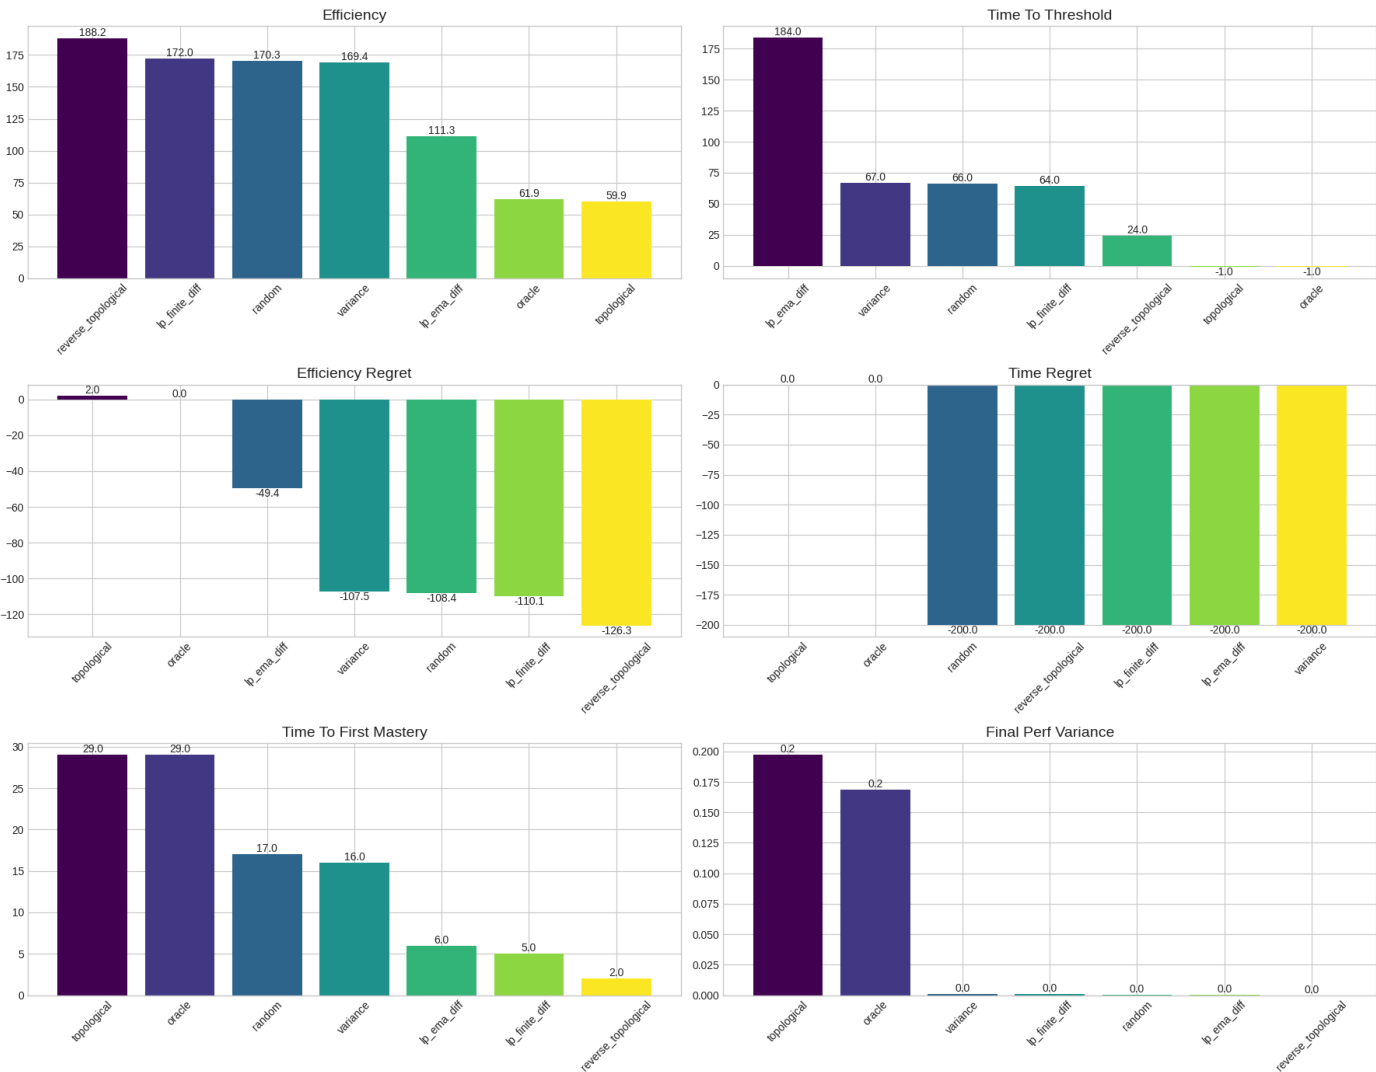
	scenario	curriculum	graph_type	efficiency	time_to_threshold	time_to_first_mastery	final_perf_1
0	1_Chain_Simple	random	chain	170.333115	66	17	6.8%
1	1_Chain_Simple	topological	chain	59.916556	-1	29	1.9%
2	1_Chain_Simple	reverse_topological	chain	188.156021	24	2	0.0%
3	1_Chain_Simple	lp_finite_diff	chain	172.032845	64	5	6.9%
4	1_Chain_Simple	lp_ema_diff	chain	111.270636	184	6	1.3%
5	1_Chain_Simple	variance	chain	169.361163	67	16	7.4%
6	1_Chain_Simple	oracle	chain	61.887785	-1	29	1.6%
7	2_Chain_HighForget	random	chain	167.824185	63	20	6.6%
8	2_Chain_HighForget	topological	chain	58.821264	-1	28	1.7%
9	2_Chain_HighForget	reverse_topological	chain	187.235618	24	2	1.1%
10	2_Chain_HighForget	lp_finite_diff	chain	172.371547	65	6	7.1%
11	2_Chain_HighForget	lp_ema_diff	chain	109.621859	185	6	2.1%
12	2_Chain_HighForget	variance	chain	169.088831	67	16	6.7%
13	2_Chain_HighForget	oracle	chain	60.436184	-1	29	1.9%
14	3_Tree_Divergent	random	tree	186.776147	55	17	7.6%
15	3_Tree_Divergent	topological	tree	38.272161	-1	200	8.6%
16	3_Tree_Divergent	reverse_topological	tree	188.741723	24	2	8.5%
17	3_Tree_Divergent	lp_finite_diff	tree	186.622626	47	13	7.1%
18	3_Tree_Divergent	lp_ema_diff	tree	177.935767	81	10	7.8%
19	3_Tree_Divergent	variance	tree	186.499868	53	16	6.7%
20	3_Tree_Divergent	oracle	tree	31.211558	-1	-1	5.5%
21	4_InvertedTree_Convergent	random	inverted_tree	177.925998	32	3	5.4%
22	4_InvertedTree_Convergent	topological	inverted_tree	188.712836	23	3	1.1%
23	4_InvertedTree_Convergent	reverse_topological	inverted_tree	188.009549	24	2	4.2%
24	4_InvertedTree_Convergent	lp_finite_diff	inverted_tree	177.705068	33	2	6.4%
25	4_InvertedTree_Convergent	lp_ema_diff	inverted_tree	169.647550	62	3	8.9%
26	4_InvertedTree_Convergent	variance	inverted_tree	176.611744	32	3	6.6%
27	4_InvertedTree_Convergent	oracle	inverted_tree	175.801671	27	4	7.2%
28	5_ComplexDAG_LowTransfer	random	dag	171.140873	79	16	7.6%
29	5_ComplexDAG_LowTransfer	topological	dag	20.379538	-1	-1	6.4%

30	5_ComplexDAG_LowTransfer	reverse_topological	dag	187.979976	24	2	6.6
31	5_ComplexDAG_LowTransfer	lp_finite_diff	dag	175.494895	65	6	9.3
32	5_ComplexDAG_LowTransfer	lp_ema_diff	dag	144.764146	138	6	1.0
33	5_ComplexDAG_LowTransfer	variance	dag	171.462537	78	17	7.5
34	5_ComplexDAG_LowTransfer	oracle	dag	42.932153	-1	43	1.5

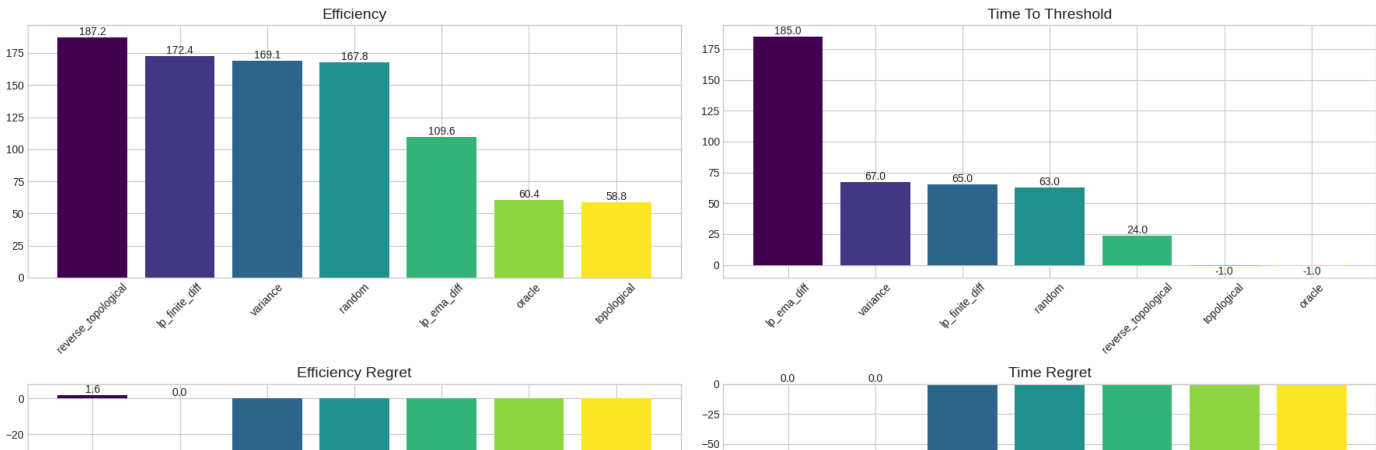
Detailed Bar Chart Comparisons

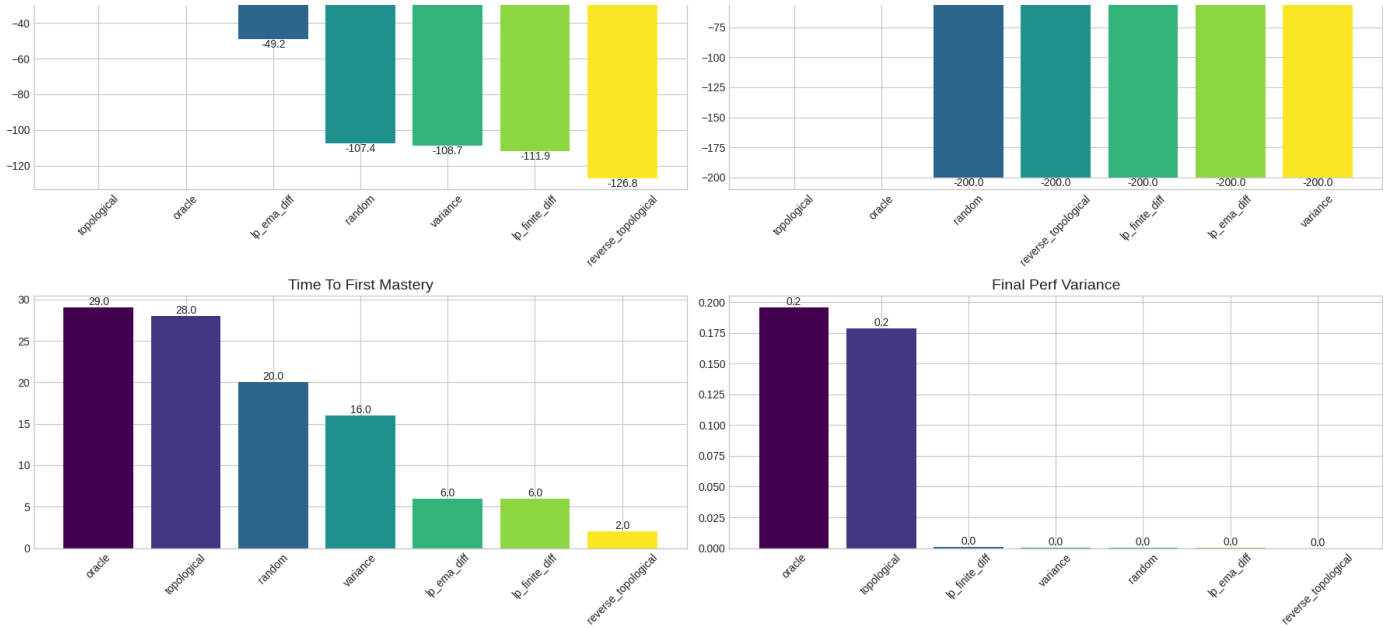
These charts provide a detailed look at curriculum performance in each specific scenario.

Metric Comparison for Scenario: 1\_Chain\_Simple

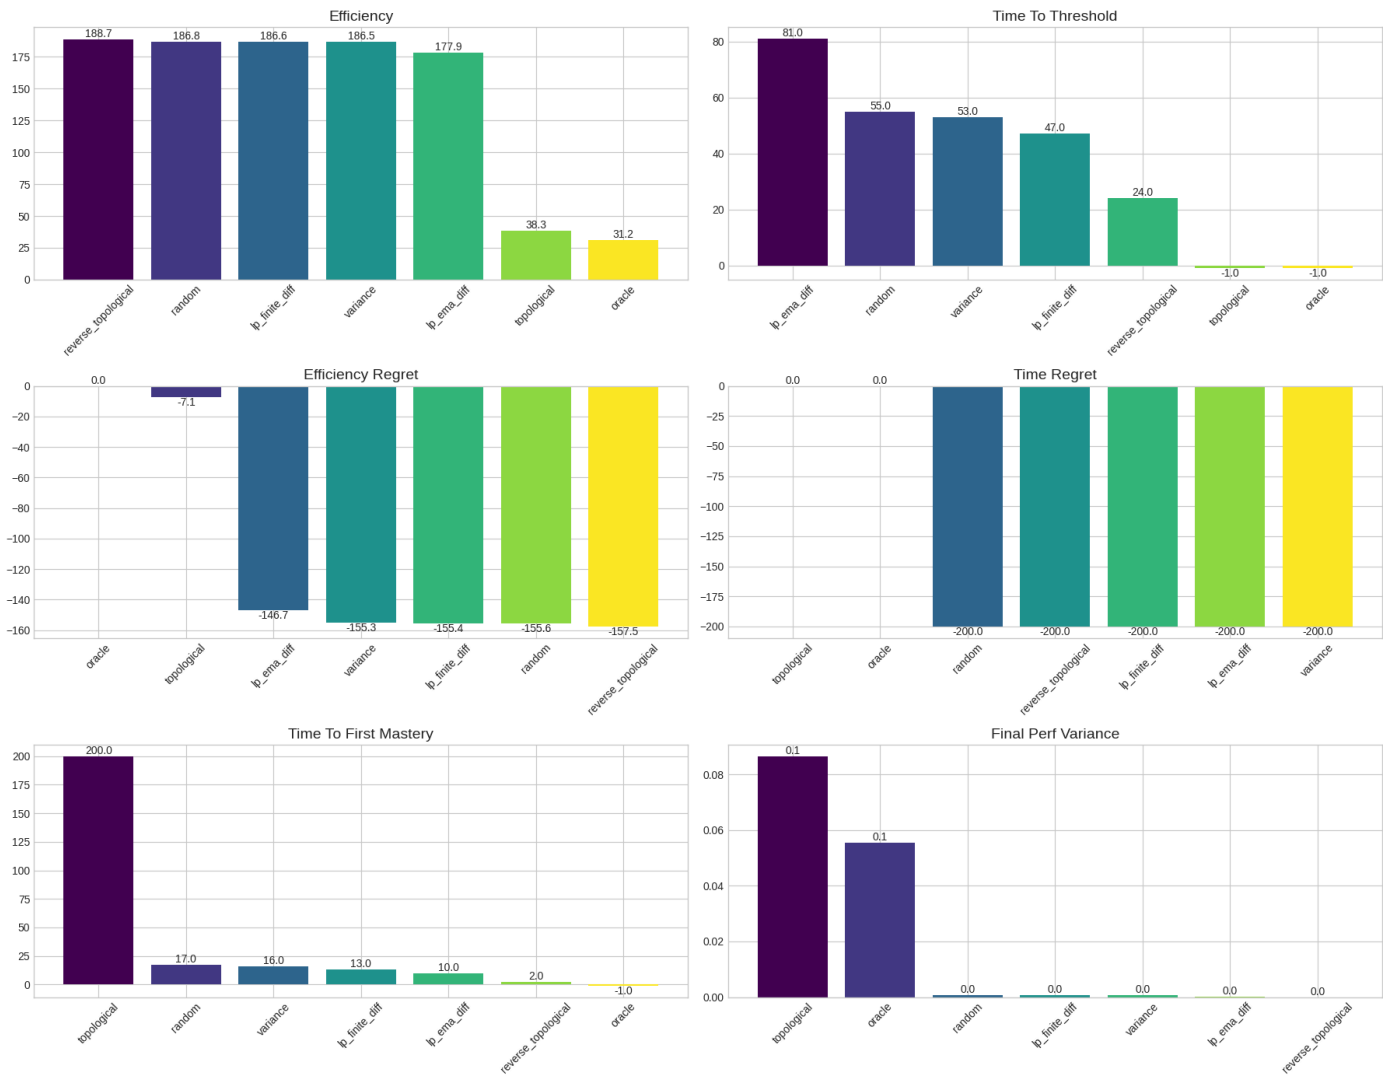


Metric Comparison for Scenario: 2\_Chain\_HighForget

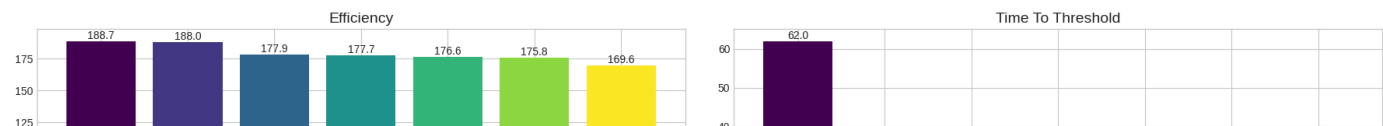


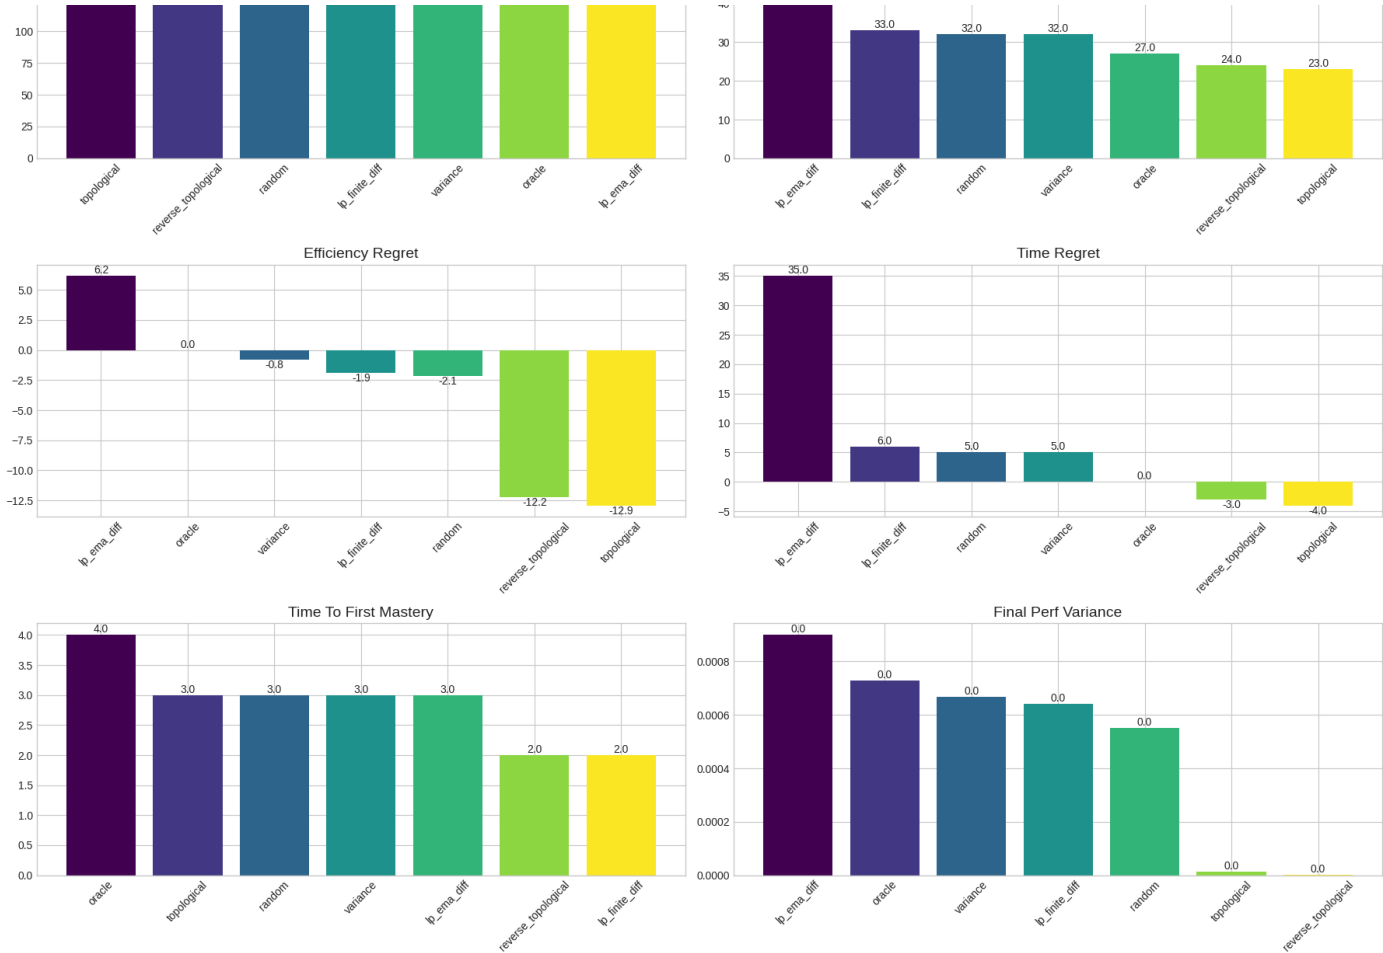


Metric Comparison for Scenario: 3\_Tree\_Divergent

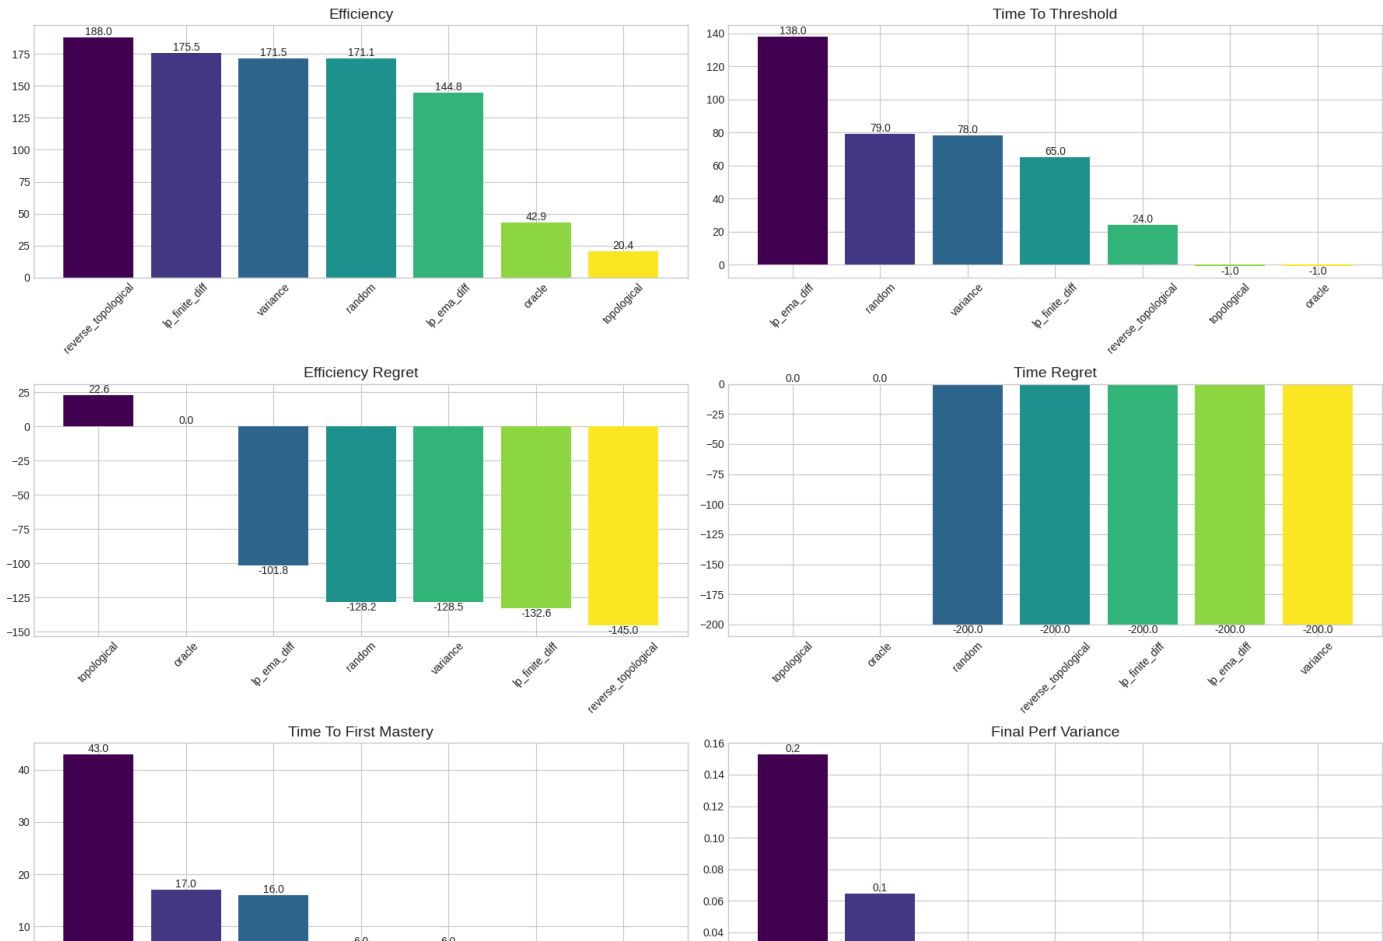


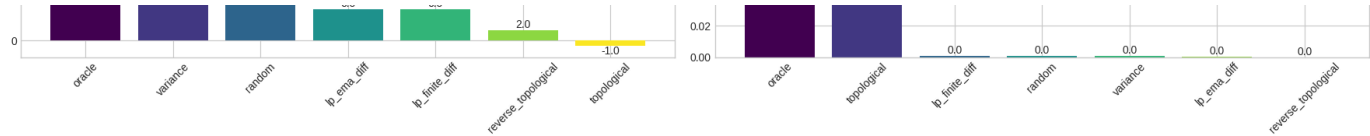
Metric Comparison for Scenario: 4\_InvertedTree\_Convergent





Metric Comparison for Scenario: 5\_ComplexDAG\_LowTransfer





Global Star Plot Comparison

This star plot is the final, high-level summary. Each point on the star represents a normalized performance metric, where a larger area is universally better. This visualization quickly reveals the unique strengths and weaknesses of each curriculum across the different environmental challenges. For example, a curriculum that excels in the 'High Forget' scenario will have a large area in that respective plot.

Curriculum Performance Profiles (Normalized)

1\_Chain\_Simple

2\_Chain\_HighForget

3\_Tree\_Divergent

